



S.B. JAIN INSTITUTE OF TECHNOLOGY MANAGEMENT & RESEARCH, NAGPUR

Practical 04

Aim: Implement process creation using the fork() system call in Linux to generate parent and child processes, retrieve their PID and PPID, and analyse system behavior by creating orphan and zombie processes.

Name: Sahil Wadekar

USN: CD24040

Semester / Year: 4th SEM / 2nd Year

Academic Session: 2025-26

Date of Performance:

Date of Submission:

❖ **Aim:** Implement process creation using the fork() system call in Linux to generate parent and child processes, retrieve their PID and PPID, and analyse system behaviour by creating orphan and zombie processes.

❖ **Tasks to be done in this Practical.**

1. Adam is working in an IT company. He has been given a task to reduce the load of a system by killing some of the processes running in the LINUX operating system. Which commands will he use to complete the given task with the help of the following operation?
 - Kill processes by name
 - Kill a process based on the process name
 - Kill a single process at a time with the given process ID
2. Write a program for process creation using C
 - Orphan Process
 - Zombie Process
3. Create the process using fork () system call.
 - Child Process creation
 - Parent process creation
 - PPID and PID

❖ **Objectives:**

- a. Implement process creation using the fork() system call to generate parent and child processes.
- b. Retrieve and display the Process ID (PID) and Parent Process ID (PPID) for both processes.
- c. Analyse system behaviour by creating and observing orphan and zombie processes.

❖ **Requirements:**

✓ **Hardware Requirements:**

- Processor: Minimum 1 GHz.
- RAM: 512 MB or higher.
- Storage: 100 MB free space.

✓ **Software Requirements:**

- Operating System: Linux/Unix-based.
- Compiler: GCC Compiler.
- Text Editor: Nano, Vim, or any preferred editor.

❖ **Theory:**

In many operating systems, the fork system call is an essential operation. The fork system call allows the creation of a new process. When a process calls the fork(), it duplicates itself, resulting in two processes running at the same time. The new process that is created is called a child process. It is a copy of the parent process. The fork system call is required for process creation and enables many important features such as parallel processing, multitasking, and the creation of complex process hierarchies.

It develops an entirely new process with a distinct execution setting. The new process has its own address space, and memory, and is a perfect duplicate of the caller process.

Basic Terminologies Used in Fork System Call in Operating System

- ❖ **Process:** In an operating system, a process is an instance of a program that is currently running. It is a separate entity with its own memory, resources, CPU, I/O hardware, and files.
- ❖ **Parent Process:** The process that uses the fork system call to start a new child process is referred to as the parent process. It acts as the parent process's beginning point and can go on running after the fork.
- ❖ **Child Process:** The newly generated process as a consequence of the fork system call is referred to as the child process. It has its own distinct process ID (PID), and memory, and is a duplicate of the parent process.
- ❖ **Process ID:** A process ID (PID) is a special identification that the operating system assigns to each process.
- ❖ **Copy-on-Write:** The fork system call makes use of the memory management strategy known as copy-on-write. Until one of them makes changes to the shared memory, it enables the parent and child processes to share the same physical memory. To preserve data integrity, a second copy is then made.
- ❖ **Return Value:** The fork system call's return value gives both the parent and child process information. It assists in handling mistakes during process formation and determining the execution route.

Process Creation

When the fork system call is used, the operating system completely copies the parent process to produce a new child process. The memory, open file descriptors, and other pertinent properties of the parent process are passed down to the child process. The child process, however, has a unique execution route and PID.

The copy-on-write method is used by the fork system call to maximize memory use. At first, the physical memory pages used by the parent and child processes are the same. To avoid unintentional changes, a separate copy is made whenever either process alters a shared memory page.

The return value of the fork call can be used by the parent to determine the execution path of the child process. If it returns 0 then it is executing the child process, if it returns -1 then there is some error; and if it returns some positive value, then it is the PID of the child process.

Advantages of Fork System Call

- **Creating new processes** with the fork system call facilitates the running of several tasks concurrently within an operating system. The system's efficiency and multitasking skills are improved by this concurrency.
- **Code reuse:** The child process inherits an exact duplicate of the parent process, including every code segment, when the fork system call is used. By using existing code, this feature encourages code reuse and streamlines the creation of complicated programmes.
- **Memory Optimisation:** When using the fork system call, the copy-on-write method optimises the use of memory. Initial memory overhead is minimised since parent and child processes share the same physical memory pages. Only when a process changes a shared memory page, improving memory efficiency, does copying take place.
- Process isolation is achieved by giving each process started by the fork system call its own memory area and set of resources. System stability and security are improved because of this isolation, which prevents processes from interfering with one another.

Disadvantages of Fork System Call

- **Memory Overhead:** The fork system call has memory overhead even with the copy-on-write optimisation. The parent process is first copied in its entirety, including all of its memory, which increases memory use.
- **Duplication of Resources:** When a process forks, the child process duplicates all open file descriptors, network connections, and other resources. This duplication may waste resources and perhaps result in inefficiencies.
- **Communication complexity:** The fork system call generates independent processes that can need to coordinate and communicate with one another. To enable data transmission across processes, interprocess communication

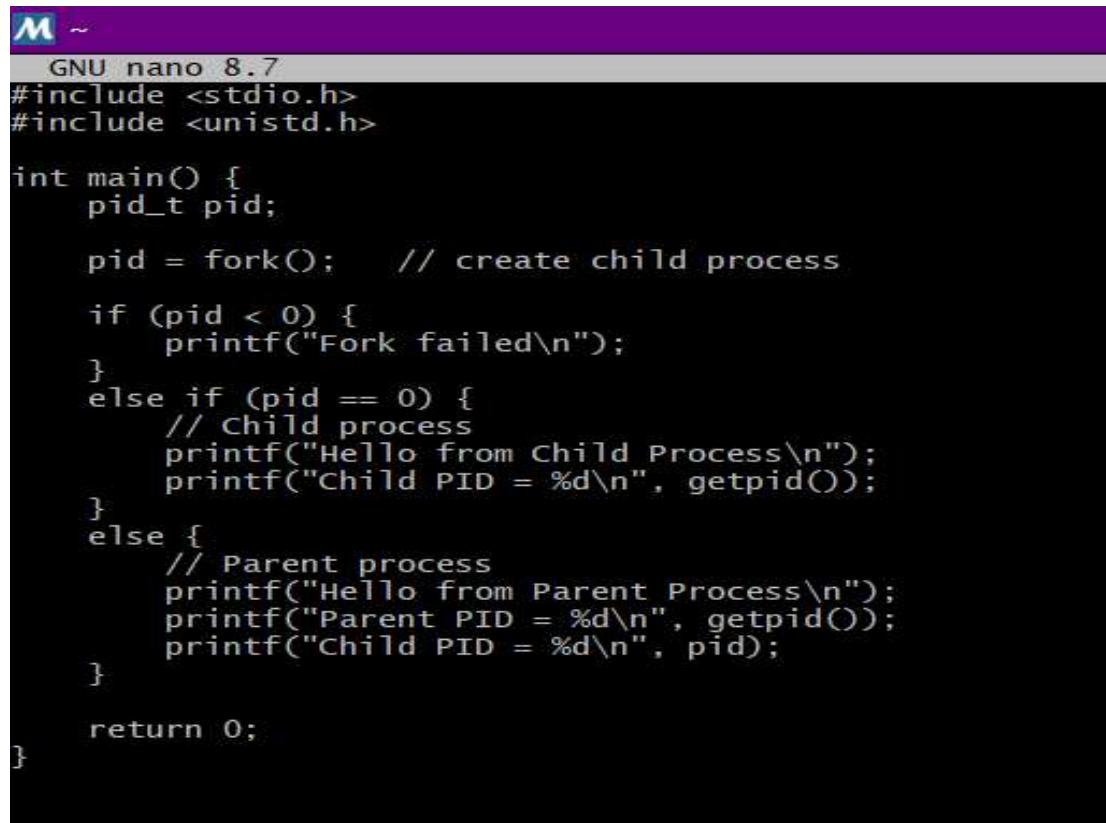
methods, such as pipes or shared memory, must be built, which might add complexity.

- **Impact on System Performance:** Forking a process duplicates memory allocation, resource management, and other system tasks. Performance of the system may be affected by this, particularly in situations where processes are often started and stopped.

❖ CODES

1. FORK.c :

```
rahul@LAPTOP-NDMNM2UM MSYS ~
$ nano fork.c
```



```
M ~
GNU nano 8.7
#include <stdio.h>
#include <unistd.h>

int main() {
    pid_t pid;

    pid = fork(); // create child process

    if (pid < 0) {
        printf("Fork failed\n");
    }
    else if (pid == 0) {
        // child process
        printf("Hello from Child Process\n");
        printf("Child PID = %d\n", getpid());
    }
    else {
        // Parent process
        printf("Hello from Parent Process\n");
        printf("Parent PID = %d\n", getpid());
        printf("Child PID = %d\n", pid);
    }

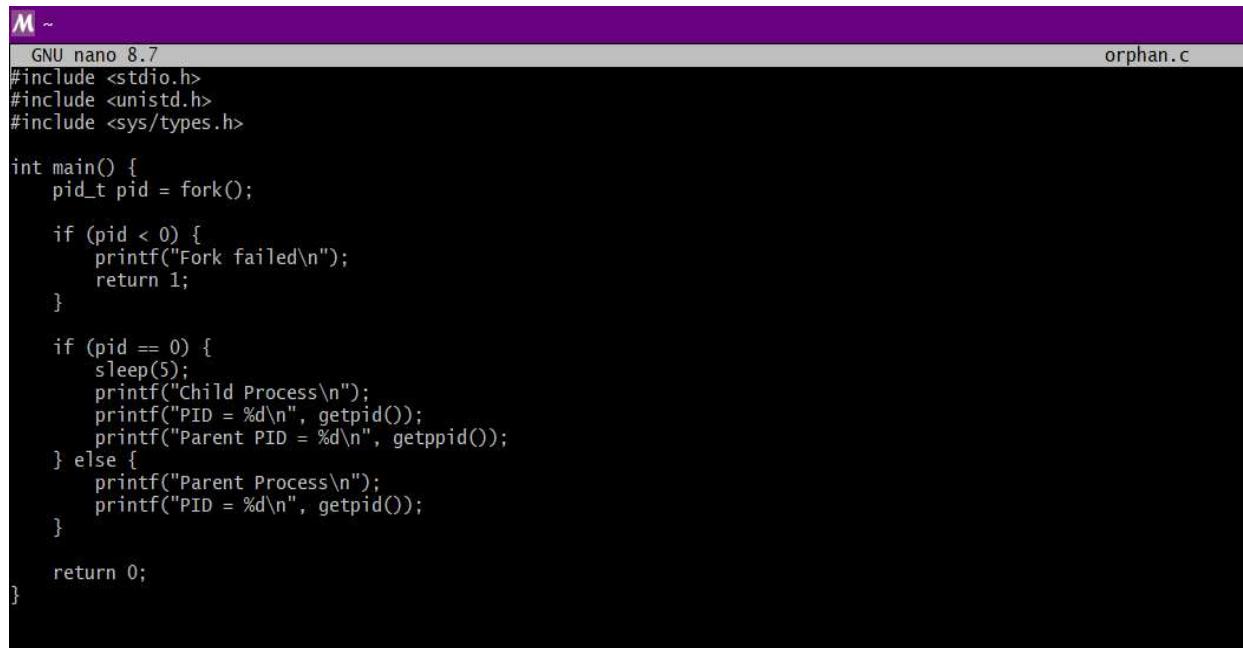
    return 0;
}
```

Output :

```
rahul@LAPTOP-NDMNM2UM MSYS ~
$ gcc fork.c -o fork

rahul@LAPTOP-NDMNM2UM MSYS ~
$ ./fork
Hello from Child Process
Hello from Parent Process
Child PID = 347
Parent PID = 346
Child PID = 347
```

2. Orphan.c:



The screenshot shows a terminal window with a purple header bar. The header bar contains the text 'M ~' on the left and 'orphan.c' on the right. The main area of the terminal is a black background with white text. It displays the C code for 'orphan.c'. The code uses the `nano` text editor's syntax highlighting, where comments are in green and strings are in red. The code itself is a simple program that forks a process. If the fork fails, it prints an error message and returns 1. If the fork succeeds, it sleeps for 5 seconds. Then, it checks if it's a child or parent process. If it's a child, it prints 'Child Process' and its PID. If it's a parent, it prints 'Parent Process' and its PID.

```
GNU nano 8.7
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

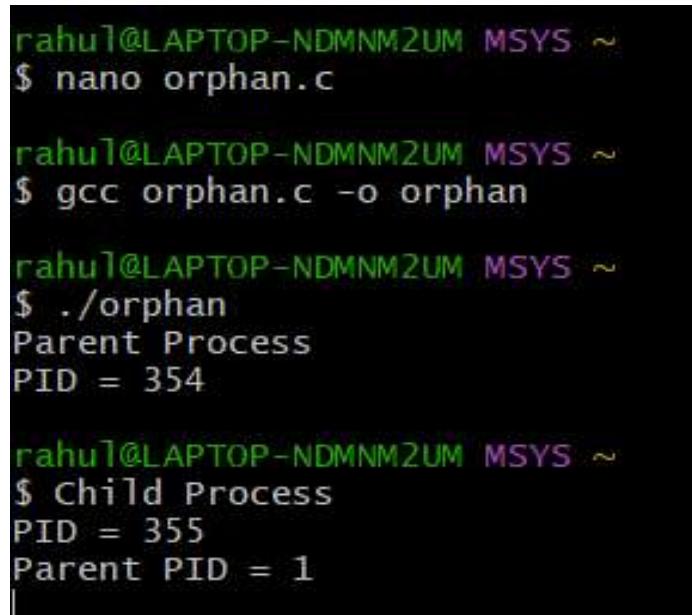
int main() {
    pid_t pid = fork();

    if (pid < 0) {
        printf("Fork failed\n");
        return 1;
    }

    if (pid == 0) {
        sleep(5);
        printf("Child Process\n");
        printf("PID = %d\n", getpid());
        printf("Parent PID = %d\n", getppid());
    } else {
        printf("Parent Process\n");
        printf("PID = %d\n", getpid());
    }

    return 0;
}
```

Output:



The screenshot shows a terminal window with a black background and white text. It displays the execution of the 'orphan.c' program. The user first runs `nano orphan.c` to view the source code. Then, they compile it with `gcc orphan.c -o orphan`. Finally, they run the compiled executable with `./orphan`. The output shows two distinct processes: a 'Parent Process' with PID 354 and a 'Child Process' with PID 355. The child process also outputs its parent's PID as 1.

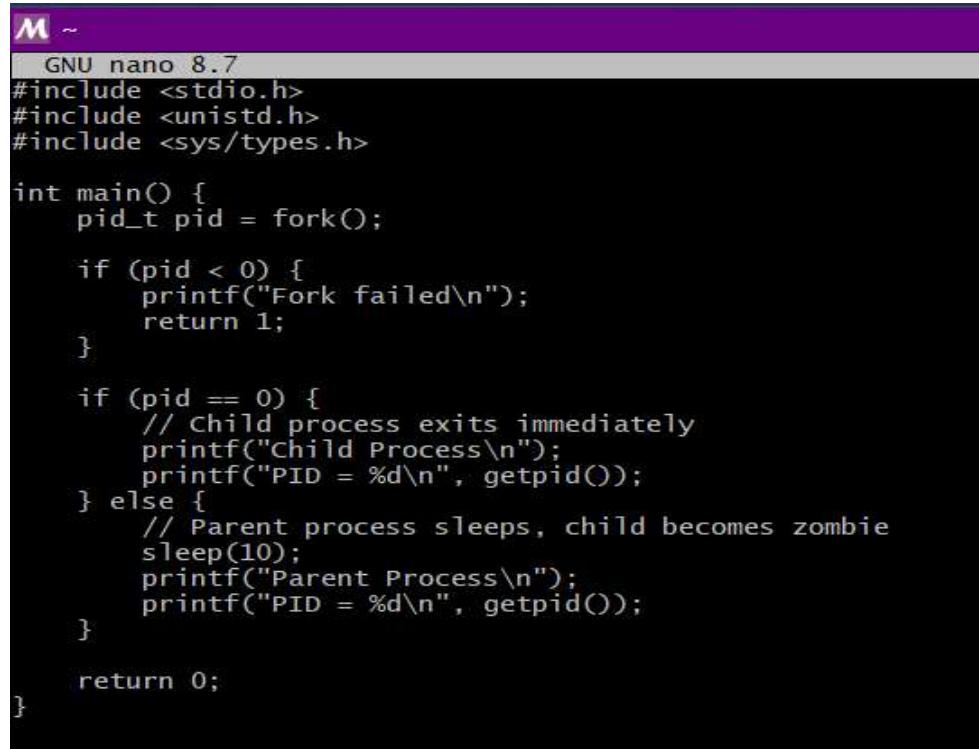
```
rahu1@LAPTOP-NDNMN2UM MSYS ~
$ nano orphan.c

rahu1@LAPTOP-NDNMN2UM MSYS ~
$ gcc orphan.c -o orphan

rahu1@LAPTOP-NDNMN2UM MSYS ~
$ ./orphan
Parent Process
PID = 354

rahu1@LAPTOP-NDNMN2UM MSYS ~
$ Child Process
PID = 355
Parent PID = 1
```

3. Zombie.c:



```
GNU nano 8.7
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

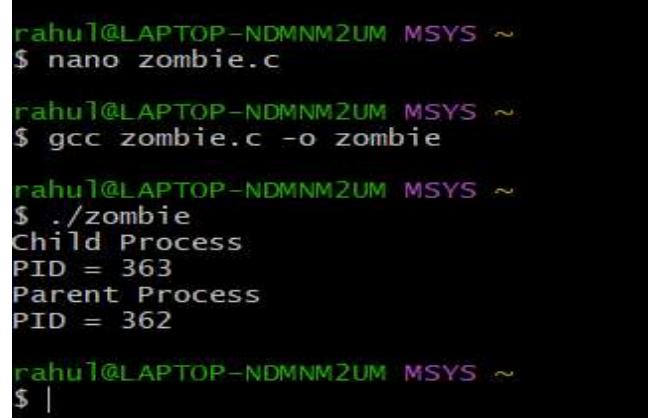
int main() {
    pid_t pid = fork();

    if (pid < 0) {
        printf("Fork failed\n");
        return 1;
    }

    if (pid == 0) {
        // Child process exits immediately
        printf("Child Process\n");
        printf("PID = %d\n", getpid());
    } else {
        // Parent process sleeps, child becomes zombie
        sleep(10);
        printf("Parent Process\n");
        printf("PID = %d\n", getpid());
    }

    return 0;
}
```

Output:



```
rahul@LAPTOP-NDMNM2UM MSYS ~
$ nano zombie.c

rahul@LAPTOP-NDMNM2UM MSYS ~
$ gcc zombie.c -o zombie

rahul@LAPTOP-NDMNM2UM MSYS ~
$ ./zombie
Child Process
PID = 363
Parent Process
PID = 362

rahul@LAPTOP-NDMNM2UM MSYS ~
$ |
```

- ❖ **Conclusion:** In this practical, we conclude that Linux commands like pkill, killall, and kill effectively reduce system load by terminating processes based on their name or ID, improving system performance and stability.

❖ **Discussion Questions:**

1. What is the concept of pointers in C?

Ans: A pointer in C is a variable that stores the memory address of another variable. It enables direct access to memory, facilitating dynamic memory management and efficient function calls, especially for large data structures.

2. What is the difference between malloc and calloc in C?

Ans: malloc allocates uninitialized memory, while calloc allocates memory and initializes all bits to zero. Use malloc when initialization is not required and calloc when zeroed memory is needed.

3. What is the purpose of the fork() system call in Unix?

Ans: fork() creates a new process by duplicating the parent process. It allows for parallel execution of processes, where the child process gets a copy of the parent's address space.

4. What is the difference between struct and union in C?

Ans: A struct allocates separate memory for each member, while a union shares the same memory space for all members, saving memory but allowing only one member to hold a value at a time.

5. What does grep do in Linux?

Ans: grep searches for specific patterns in files and outputs the lines containing those patterns. It's a powerful tool for text processing and log analysis in Linux environments.

❖ **References:**

https://www.geeksforgeeks.org/fork-system-call-in-operating-system/?ref=ml_lbp

<https://www.scaler.com/topics/fork-system-call/>

Date: _____ / _____ /2026

Signature

Course Coordinator

B.Tech CSE(DS)

Sem: 4 / 2025-26

Operating System Lab (N-PCCCD401P)